

Assignment 2

Part(A):

1-

Reading the data

```
1 dataset <- read.csv(file = 'C:/Users/Genius/Downloads/Assignment 2/Churn Dataset.csv')
2 head(dataset)
3 df=as.data.frame(do.call(cbind, dataset))
4 head(df)
5
6
```

This is the result

```
> head(df)
  customerID gender SeniorCitizen Partner Dependents tenure PhoneService MultipleLines
1 7590-VHVEG Female            0      Yes         No         1           No No phone service
2 5575-GNVDE Male              0      No         No        34           Yes           No
3 3668-QPYBK Male              0      No         No        45           Yes           No
4 7795-CFOCW Male              0      No         No         2           No No phone service
5 9237-HQITU Female            0      No         No         2           Yes           No
6 9305-CDSKC Female            0      No         No         8           Yes           Yes
  InternetService OnlineSecurity OnlineBackup DeviceProtection TechSupport StreamingTV StreamingMovies
1 DSL No Yes No No No No No
2 DSL Yes No No No No No No
3 DSL Yes Yes No No No No No
4 DSL Yes No No No Yes No No
5 Fiber optic No No No No No No No
6 Fiber optic No No No Yes No No Yes
  Contract PaperlessBilling PaymentMethod MonthlyCharges TotalCharges Churn
1 Month-to-month Yes Electronic check 29.85 29.85 No
2 One year No Mailed check 56.95 1889.5 No
3 Month-to-month Yes Mailed check 53.85 108.15 Yes
4 One year No Bank transfer (automatic) 42.3 1840.75 No
5 Month-to-month Yes Electronic check 70.7 151.65 Yes
6 Month-to-month Yes Electronic check 99.65 820.5 Yes
```

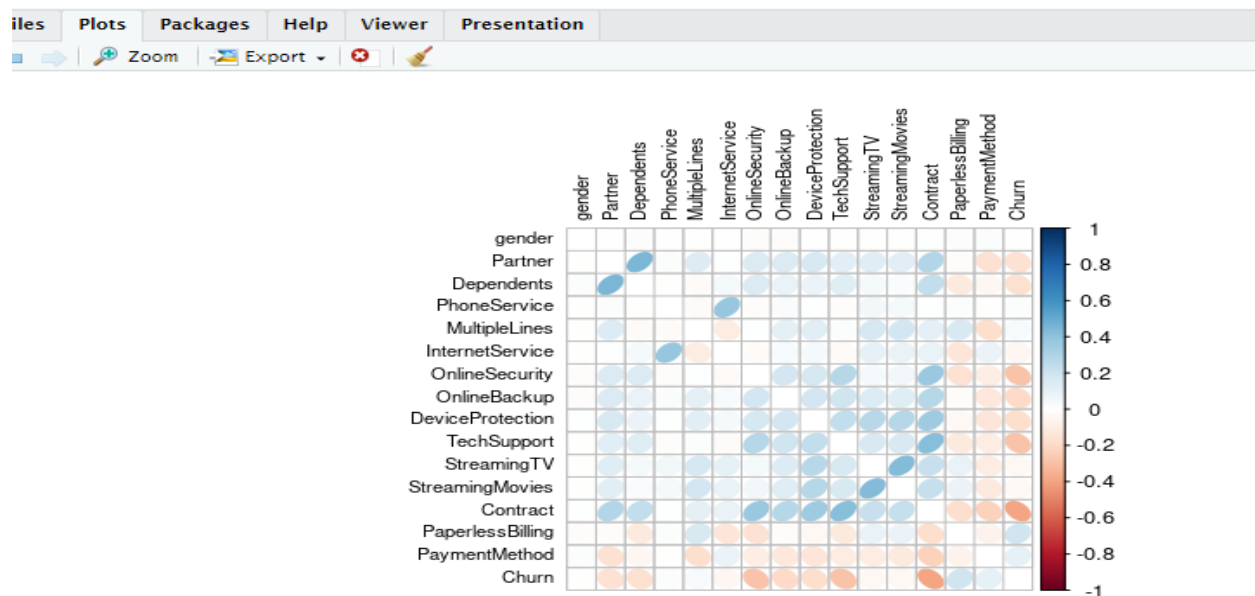
Scatter blot matrix

```
29 install.packages("corrplot")
30 library(corrplot)
31 library(reshape2)
32 library(ggplot2)
33 print((df[c(3,6,19,20)]))
34
35 corrplot(cor(df[, sapply(df, is.numeric)],method = "pearson"),diag = FALSE,
36          method = "ellipse",
37          tl.cex = 0.7, tl.col = "black", cl.ratio = 0.2
38 )|
```

And this is the scatter blot matrix of the numerical columns in the churn dataset



And this is scatter blot matrix to all data after convert categorical values to numerical



2-Heatmap

```

23 #-----
24 #load reshape2 package to use melt() function
25 library(reshape2)
26
27 #melt data into long format
28 melt_churnDataSet<- melt(cor(dataset[, sapply(dataset, is.numeric)],
29                             method = "pearson"))
30 #create heatmap using rescaled values
31 ggplot(melt_churnDataSet, aes(var1, var2)) +
32   geom_tile(aes(fill = value), colour = "white") +
33   scale_fill_gradient(low = "white", high = "red")
34

```

And that the heatmap to determine correlated attributes



Convert categorical to numerical values

```

16
17 df$gender <- as.numeric(as.factor(df$gender))
18 df$Partner <- as.numeric(as.factor(df$Partner))
19 df$Dependents <- as.numeric(as.factor(df$Dependents))
20 df$PhoneService <- as.numeric(as.factor(df$PhoneService))
21 df$MultipleLines <- as.numeric(as.factor(df$MultipleLines))
22 df$InternetService <- as.numeric(as.factor(df$InternetService))
23 df$OnlineSecurity <- as.numeric(as.factor(df$OnlineSecurity))
24 df$OnlineBackup <- as.numeric(as.factor(df$OnlineBackup))
25 df$DeviceProtection <- as.numeric(as.factor(df$DeviceProtection))
26 df$TechSupport <- as.numeric(as.factor(df$TechSupport))
27 df$StreamingTV <- as.numeric(as.factor(df$StreamingTV))
28 df$StreamingMovies <- as.numeric(as.factor(df$StreamingMovies))
29 df$Contract <- as.numeric(as.factor(df$Contract))
30 df$PaperlessBilling <- as.numeric(as.factor(df$PaperlessBilling))
31 df$Contract <- as.numeric(as.factor(df$Contract))
32 df$PaymentMethod <- as.numeric(as.factor(df$PaymentMethod))
33 df$Churn <- as.numeric(as.factor(df$Churn))
34 head(df)
35

```

And this the data after convert categorical to numerical

```

> head(df)
  CustomerID gender SeniorCitizen Partner Dependents tenure PhoneService MultipleLines
1 7590-VHVEG      1             0       2           1         1           1           2
2 5575-GNVDE      2             0       1           1         34           2           1
3 3668-QPYBK      2             0       1           1         2           2           1
4 7795-CFOCW      2             0       1           1         45           1           2
5 9237-HQITU      1             0       1           1         2           2           1
6 9305-CDSKC      1             0       1           1         8           2           3
  InternetService OnlineSecurity OnlineBackup DeviceProtection TechSupport StreamingTV
1               1              1             3                1           1           1
2               1              1             1                3           1           1
3               1              3             3                1           1           1
4               1              3             1                3           3           1
5               2              1             1                1           1           1
6               2              1             1                3           1           3
  StreamingMovies Contract PaperlessBilling PaymentMethod MonthlyCharges TotalCharges Churn
1               1         1               2             3         29.85         29.85      1
2               1         2               1             4         56.95        1889.5      1
3               1         1               2             4         53.85        108.15      2
4               1         2               1             1         42.3         1840.75     1
5               1         1               2             3         70.7         151.65     2
6               3         1               2             3         99.65         820.5      2

```

Then I will remove the customerID and TotalChrgae features from the data

```
57  
58 df<-df %>% select(-customerID)  
59 df<-df %>% select(-TotalCharges)  
60  
61 |
```

Check if there is a missing values in the data

```
1  
2 sum(is.na(df))  
3 colSums(is.na(df))  
4
```

Give me this result, there is no missing values

```
> sum(is.na(df))  
[1] 0  
> colSums(is.na(df))  
      gender SeniorCitizen      Partner      Dependents      tenure PhoneService MultipleLines  
      0         0           0         0           0         0           0  
InternetService OnlineSecurity OnlineBackup DeviceProtection TechSupport StreamingTV StreamingMovies  
      0         0           0         0           0         0           0  
      Contract PaperlessBilling PaymentMethod MonthlyCharges      Churn  
      0         0           0         0           0
```

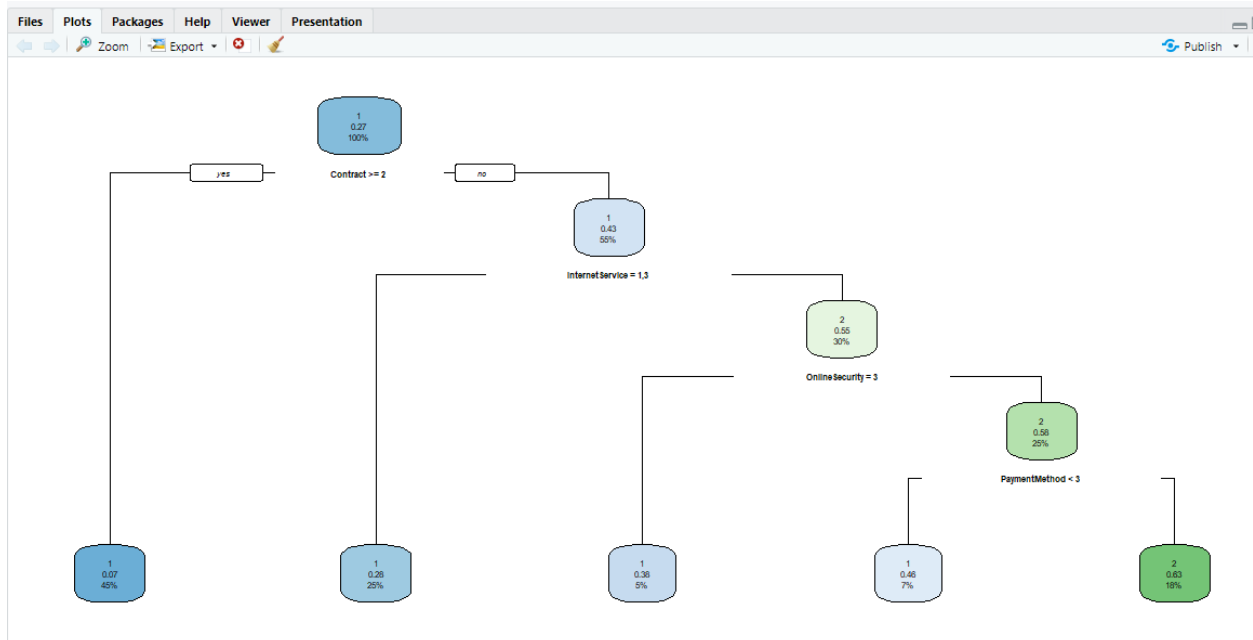
3-First Install some libraries

```
68  
69 install.packages("caTools")  
70 library(caTools)  
71 install.packages("rpart")  
72 library(rpart)  
73 install.packages("rpart.plot")  
74 library(rpart.plot)  
75 install.packages("caret")  
76 library(caret)  
77 install.packages("dplyr")  
78 library(dplyr)  
79 install.packages("lattice")  
80 library(lattice)  
81 install.packages("party")  
82 library(party)  
83 |
```

Then split the data to train set and test set

```
84  
85 set.seed(42)  
86 sample_split <- sample.split(Y = df$Churn, SplitRatio = 0.80)  
87 train_set <- subset(x = df, sample_split == TRUE)  
88 test_set <- subset(x = df, sample_split == FALSE)  
89 train_set  
90 model <- rpart(Churn ~ ., data = train_set, method = "class") #specify method as class since we are dealing with  
91 model  
92 #plot the model  
93 rpart.plot(model)  
94
```

This is the tree



```

122 #Make predictions
123 preds <- predict(model, newdata = test_set, type = "class") #use the predict() function and pass in t
124 preds
125
126 #Print the confusion Matrix
127 confusionMatrix(test_set$Churn, preds) #check accuracy
128

```

And this is the results getting accuracy=77.5%

```

> confusionMatrix(test_set$Churn, preds) #check accuracy
Confusion Matrix and Statistics

      Reference
Prediction  1    2
 1  946   89
 2  228  146

      Accuracy : 0.775
      95% CI   : (0.7523, 0.7966)
 No Information Rate : 0.8332
 P-value [Acc > NIR] : 1

      Kappa : 0.3454

 Mcnemar's Test P-value : 9.128e-15

      Sensitivity : 0.8058
      Specificity : 0.6213
   Pos Pred Value : 0.9140
   Neg Pred Value : 0.3904
     Prevalence : 0.8332
 Detection Rate : 0.6714
Detection Prevalence : 0.7346
 Balanced Accuracy : 0.7135

      'Positive' Class : 1
> |

```

4-Try different ways to improve the decision tree algorithm (e.g., use different splitting strategies, prune tree after splitting). Does pruning the tree improves the accuracy?

First splitting by Information gain

```
decisionTreeInformation <- rpart(Churn ~ ., data = train_set, method = "class", parms = list(split = "information")) #specify m
decisionTreeInformation
```

```
14 #Make predictions
15 preds <- predict(decisionTreeInformation, newdata = test_set, type = "class") #use the predict() function and pass in t
16 preds
17 #Print the confusion Matrix
18 confusionMatrix(test_set$Churn, preds) #check accuracy
19
20
21
```

Getting accuracy =77.5% , and this is the Confusion matrix

```
> confusionMatrix(test_set$Churn, preds) #check accuracy
Confusion Matrix and Statistics

          Reference
Prediction  1    2
          1  946  89
          2  228 146

              Accuracy : 0.775
              95% CI   : (0.7523, 0.7966)
              No Information Rate : 0.8332
              P-Value [Acc > NIR] : 1

              Kappa : 0.3454

  Mcnemar's Test P-Value : 9.128e-15

              Sensitivity : 0.8058
              Specificity : 0.6213
              Pos Pred Value : 0.9140
              Neg Pred Value : 0.3904
              Prevalence : 0.8332
              Detection Rate : 0.6714
              Detection Prevalence : 0.7346
              Balanced Accuracy : 0.7135

              'Positive' Class : 1
```

Second splitting the data by Gini Index

```
149 |
150 #Splitting by Gini Index
151 decisionTreeGini <- train(Churn ~ ., data = test_set,
152                           method = "rpart",
153                           parms = list(split = "gini"),
154                           trControl = ct,
155                           tuneLength = 100)
156 #Make predictions
157 preds <- predict(decisionTreeGini, newdata = test_set, type = "class") #use the predict() function and pass in the
158 preds
159 #Print the confusion Matrix
160 confusionMatrix(test_set$Churn, preds) #check accuracy
161
```

Getting accuracy =77.5% after splitting by gini index , and this is the Confusion matrix

```
Console Terminal Jobs
R 4.2.1 - /cloud/project/
> confusionMatrix(test_set$Churn, preds) #check accuracy
Confusion Matrix and Statistics

      Reference
Prediction 1 2
1  931 104
2  212 162

      Accuracy : 0.7757
      95% CI : (0.753, 0.7973)
      No Information Rate : 0.8112
      P-Value [Acc > NIR] : 0.9996

      Kappa : 0.3665

      Mcnemar's Test P-Value : 1.753e-09

      Sensitivity : 0.8145
      Specificity : 0.6090
      Pos Pred Value : 0.8995
      Neg Pred Value : 0.4332
      Prevalence : 0.8112
      Detection Rate : 0.6608
      Detection Prevalence : 0.7346
      Balanced Accuracy : 0.7118

      'Positive' Class : 1
```

Then Puruning the tree

```
Untitled1*
163 #pruning
164 decitionTreePrune <- rpart(Churn ~ ., data = train_set, method = "class",
165                           control = rpart.control(cp = 0.0082, maxdepth = 3, minsplit = 2))
166 preds<- predict(decitionTreePrune, newdata = test_set, type = "class") #use the predict() function and pass in the testing subset
167 preds
168 #Print the confusion Matrix
169 confusionMatrix(test_set$Churn, preds) #check accuracy
170 <

171:1 (Top Level)
R 4.2.1 - /cloud/project/
[reached getOption("max.print") -- omitted 409 entries ]
Levels: 1 2
> confusionMatrix(test_set$Churn, preds) #check accuracy
```

Getting accuracy =77.7% after puruning , and this is the Confusion matrix

```
> confusionMatrix(test_set$Churn, preds) #check accuracy
Confusion Matrix and Statistics

      Reference
Prediction 1 2
1  926 109
2  218 156

      Accuracy : 0.7679
      95% CI : (0.745, 0.7897)
      No Information Rate : 0.8119
      P-Value [Acc > NIR] : 1

      Kappa : 0.3438

      Mcnemar's Test P-Value : 2.338e-09

      Sensitivity : 0.8094
      Specificity : 0.5887
      Pos Pred Value : 0.8947
      Neg Pred Value : 0.4171
      Prevalence : 0.8119
      Detection Rate : 0.6572
      Detection Prevalence : 0.7346
      Balanced Accuracy : 0.6991

      'Positive' Class : 1
```

Does pruning the tree improves the accuracy? Yes To some extent

5-first installing the libraries then split the train data to (X_train , y_train) and test data to (X_test ,y_test) and create a` xgboost model to train the data and and get the accuracy=78%

```
144 install.packages('xgboost')      # for fitting the xgboost model
145 install.packages('caret')        # for general data preparation and model fitting
146 install.packages('e1071')
147 library(xgboost)
148 library(caret)
149 library(e1071)
150 my_data=as.list(df)
151
152 X_train = train_set[,-19]          # independent variables for train
153 X_train <- matrix(unlist(X_train), ncol = 18, nrow = 5634)
154 y_train = train_set[,19]           # dependent variables for train
155
156
157 X_test = test_set[,-19]            # independent variables for test
158 X_test <- matrix(unlist(X_test), ncol = 18, nrow = 1409)
159 y_test = test_set[,19]
160
161 typeof(df)
162 # convert the train and test data into xgboost matrix type.
163 xgboost_train = xgb.DMatrix(data=X_train, label=y_train)
164 xgboost_test = xgb.DMatrix(data=X_test, label=y_test)
165 #Step 4 - Create a xgboost model
166 # train a model using our training data
167 model <- xgboost(data = xgboost_train,          # the data
168                  max.depth=3, ,                # max depth
169                  nrounds=70)                   # max number of boosting iterations
170
171 summary(model)
172 |
172:1 (Top Level) <
```

```
174
175 #use model to make predictions on test data
176 pred_test = predict(model, xgboost_test)
177 pred_test
178
179 #Step 6 - Convert prediction to factor type
180 pred_test[(pred_test>3)] = 3
181 pred_y = as.factor((levels(y_test))[round(pred_test)])
182 print(pred_y)
183 |
184 #Step 7 - Create a confusion matrix
185 conf_mat = confusionMatrix(y_test, pred_y)
186 print(conf_mat)
```

And this is the confusion matrix


```
177:16 (Untitled)
Console Terminal Background Jobs
R 4.2.1 · ~/assignment2/
> conf_mat=confusionMatrix(y_test,pred_y)
> print(conf_mat)
Confusion Matrix and Statistics

      Reference
Prediction 1  2
1  917 118
2  188 186

      Accuracy : 0.7828
      95% CI : (0.7604, 0.8041)
      No Information Rate : 0.7842
      P-Value [Acc > NIR] : 0.5667

      Kappa : 0.4077

      Mcnemar's Test P-value : 7.998e-05

      Sensitivity : 0.8299
      Specificity : 0.6118
      Pos Pred Value : 0.8860
      Neg Pred Value : 0.4973
      Prevalence : 0.7842
      Detection Rate : 0.6508
      Detection Prevalence : 0.7346
      Balanced Accuracy : 0.7209

      'Positive' class : 1

> |
```

6-Train a deep neural network using Keras with 3 dense layers. Try changing the activation function or dropout rate.

First installing keras

```
181 #
182 install.packages('devtools')
183 devtools::install_github("rstudio/keras")
184 devtools::install_github("rstudio/reticulate")
185
186 install.packages("tensorflow")
187
188 install.packages('reticulate')
189 install.packages('keras')
190 library(reticulate)
191 library(keras)
192 library(tensorflow)
193
194
195 #reticulate::use_python("C:/Users/Genius/AppData/Local/Microsoft/windowsApps/PythonSoftwareFoundation.Python.3.9
196 reticulate::use_python("C:/Users/Genius/Anaconda3/")
197 install_tensorflow()
198
```

then creating a model has 3 dense layers and fit the training data

```

200 #defining a keras sequential model
201 model <- keras_model_sequential()
202
203 model %>%
204   layer_dense(units = 19, input_shape = 784) %>%
205   layer_dropout(rate=0.5)%>%
206   layer_activation(activation = 'relu') %>%
207   layer_dense(units = 50) %>%
208   layer_activation(activation = 'softmax')
209   layer_dense(units = 2) %>%
210   layer_activation(activation = 'softmax')
211
212
213 #compiling the defined model with metric = accuracy and optimiser as adam.
214 model %>% compile(
215   loss = 'categorical_crossentropy',
216   optimizer = 'adam',
217   metrics = c('accuracy')
218 )
219
220
221 #fitting the model on the training dataset
222 model %>% fit(X_train , y_train , epochs = 50, batch_size = 128)
223
224 #Evaluating model on the cross validation dataset
225 loss_and_metrics <- model %>% evaluate(X_test , y_test, batch_size = 128)
226
227
228
229
230
231
232
233
234
235
236 preds_DNN <- predict(model , newdata = test_set, type = "class") #use the predict() function and pass in the testing subset
237 preds_DNN
238
239 #Print the confusion Matrix
240 confusionMatrix(test_set$churn, preds_DNN) #check accuracy
241

```

And this is the confusion matrix , getting accuracy=77.4%

```

44/44 [=====] - 0s 2ms/step
Confusion Matrix and Statistics

      Prediction Reference
      0      1
0  999    34
1  283    88

      Accuracy : 0.7742
      95% CI   : (0.7514, 0.7959)
No Information Rate : 0.9131
P-Value [Acc > NIR] : 1

      Kappa : 0.2603

McNemar's Test P-Value : <2e-16

      Sensitivity : 0.77793
      Specificity : 0.7213
      Pos Pred Value : 0.9671
      Neg Pred Value : 0.2372
      Precision : 0.9671
      Recall : 0.77793
      F1 : 0.8631
      Prevalence : 0.9131
      Detection Rate : 0.7115
      Detection Prevalence : 0.7358
      Balanced Accuracy : 0.7503

      'Positive' Class : 0

```

Is there any sign of overfitting? From confusion matrix ,Sensitivity (also known as recall) and Pos Pred Value(also known as precision). Then F1 can be easily computed, as stated above, as: $F1 \leftarrow (2 * \text{precision} * \text{recall}) / (\text{precision} + \text{recall})$ so $F1=0.8570$,**So there is no Overfitting**

7- Compare the performance of the models in terms of the following criteria: precision, recall, accuracy, F measure. Identify the model that performed best and worst according to each criteria

From confusion matrix ,Sensitivity (also known as recall) and Pos Pred Value(also known as precision). Then F1 can be easily computed, as stated above, as: $F1 <- (2 * precision * recall) / (precision + recall)$

Model	Accuracy	Recall	Precision	F1
Decision Tree	0.775	0.8058	0.9140	0.8564
XGboost	0.7828	0.8299	0.8860	0.8570
DNN	0.7742	0.7793	0.9671	0.8631

-The best model that perform highest Accuracy is **XGboost** ,the worst model is **DNN**

-The best model that perform highest Recall is **XGboost** ,the worst model is **DNN**

-The best model that perform highest Precision is **DNN** ,the worst model is **XGboost**

-The best model that perform highest F1 score is **DNN** ,the worst model is **Decision Tree**

8- first ROC for Decision tree

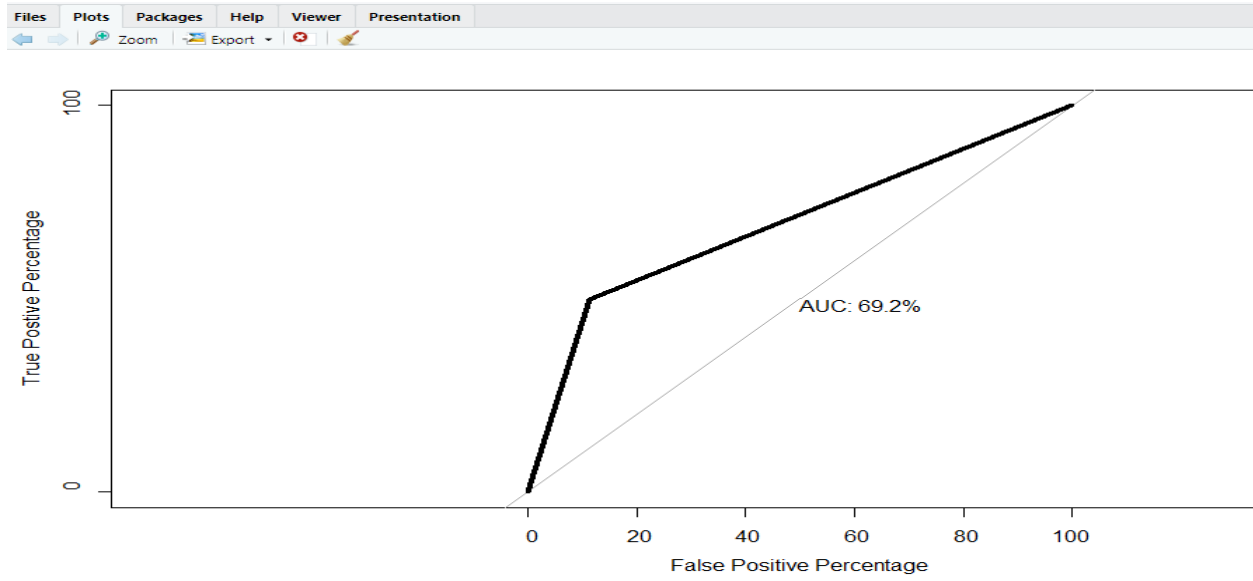


Second For ROC XGboost

```

208 roc(test_set$churn, as.numeric(as.character(pred_y)), plot=TRUE, legacy.axes=TRUE, percent=TRUE, xlab="False Positive Percentage", ylab="True Post
209 legend("bottomright", legend=c("model_name"), col=c("#000000"), lwd=4)
210

```

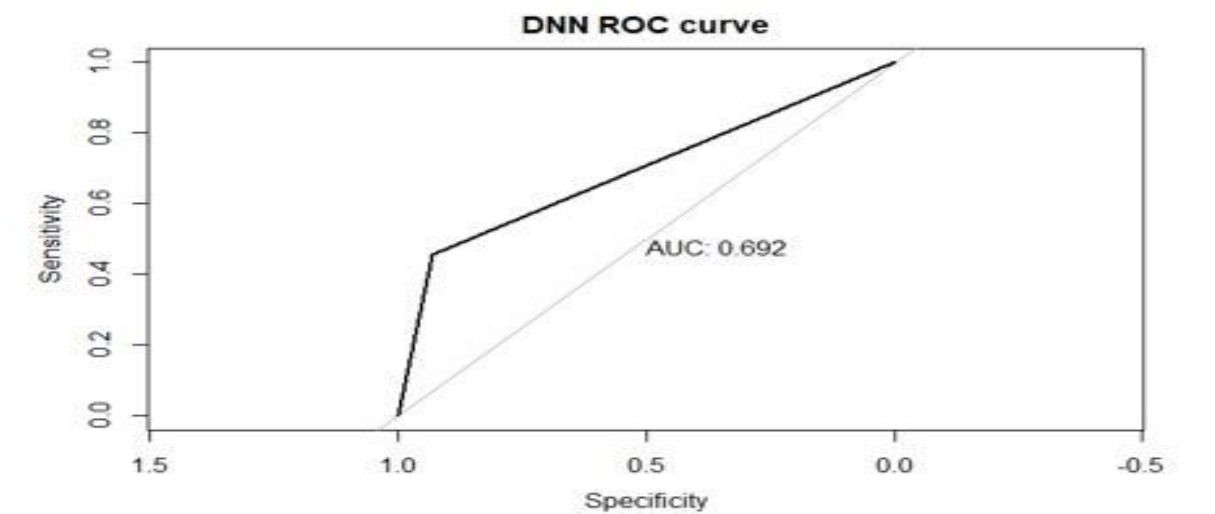


Third For ROC DNN

```

?61 #ROC graph For DNN
?62 roc(test_set$churn, as.numeric(as.character(preds_DNN)), plot=TRUE, legacy.axes=TRUE, percent=TRUE, xlab="False Positive Percentage", ylab="True Postive Pe
?63 legend("bottomright", legend=c("model_name"), col=c("#000000"), lwd=4)
?64

```



Part(B)

(a)-First reading the transactions data set

```

235 #-----
236 #Part (B) in the Assignment
237 #first reading the transactions
238
239 install.packages("arules")
240 library(arules)
241 install.packages("arulesviz")
242 library(arulesviz)
243 install.packages("readr")
244 library(readr)
245 install.packages("RColorBrewer")
246 library(RColorBrewer)
247
248 d <- read.transactions('C:/Users/Genius/Downloads/Assignment 2/transactions.csv', format = 'basket', sep = ',')
249 head(d)
250 typeof(d)
251 summary(d)
252 plot(head(d,10))

```

And this is the summary of the data

```

> summary(d)
transactions as itemMatrix in sparse format with
 7501 rows (elements/itemsets/transactions) and
 119 columns (items) and a density of 0.03288973

most frequent items:
mineral water      eggs      spaghetti french fries      chocolate      (other)
      1788          1348          1306          1282          1229          22405


element (itemset/transaction) length distribution:
sizes
 1    2    3    4    5    6    7    8    9   10   11   12   13   14   15   16   18   19   20
1754 1358 1044  816  667  493  391  324  259  139  102  67  40  22  17  4  1  2  1

  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 1.000  2.000   3.000   3.914  5.000  20.000

includes extended item information - examples:
      labels
1      almonds
2 antioxidant juice
3      asparagus
> |

```

And this is a plot of the top 10 transactions



(b)-

Generate association rules using minimum support of 0.002, minimum confidence of 0.20, and maximum length of 3

```

287 # set better support and confidence levels to learn more rules
288 transaction_rules <- apriori(d, parameter = list(support = 0.002, confidence = 0.20, maxlen = 3))
289 summary(transaction_rules)
290 plot(head(d,10))

```

Display the rules: this is a summary to rules

```

> summary(transaction_rules)
set of 2023 rules

rule length distribution (lhs + rhs): sizes
  1      2      3
  1 357 1665

  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 1.000  3.000   3.000   2.823  3.000   3.000

summary of quality measures:
  support      confidence      coverage      lift      count
Min. :0.002133  Min. :0.2000  Min. :0.002666  Min. : 0.8595  Min. : 16.0
1st Qu.:0.002533 1st Qu.:0.2405  1st Qu.:0.008266 1st Qu.: 1.5377 1st Qu.: 19.0
Median :0.003466 Median :0.2941  Median :0.011465 Median : 1.8674 Median : 26.0
Mean   :0.005292 Mean   :0.3177  Mean   :0.018647 Mean   : 2.0415 Mean   : 39.7
3rd Qu.:0.005599 3rd Qu.:0.3774  3rd Qu.:0.019064 3rd Qu.: 2.3381 3rd Qu.: 42.0
Max.   :0.238368 Max.   :0.9500  Max.   :1.000000 Max.   :28.0881 Max.   :1788.0

mining info:
data ntransactions support confidence                                call
d      7501      0.002      0.2 apriori(data = d, parameter = list(support = 0.002, confidence = 0.2, maxlen = 3))
> |

```

And this is the first two rules

```

> inspect(transaction_rules[1:2])
  lhs      rhs      support      confidence coverage      lift      count
[1] {}      => {mineral water} 0.238368218 0.2383682 1.00000000 1.000000 1788
[2] {asparagus} => {mineral water} 0.002133049 0.4444444 0.00479936 1.864529 16
> |

```

Sorting the transactions rules by descending lift value

```

294
295 # sorting descending transactions rules by lift to determine actionable rules
296 top.lift <- sort(transaction_rules, decreasing = TRUE, na.last = NA, by = "lift")
297 inspect(top.lift[1:5])
298

> top.lift1 <- sort(transaction_rules, decreasing = TRUE, na.last = NA, by = "lift")
> inspect(top.lift1[1:5])
  lhs      rhs      support      confidence coverage      lift      count
[1] {escalope, mushroom cream sauce} => {pasta} 0.002532996 0.4418605 0.005732569 28.088096 19
[2] {escalope, pasta} => {mushroom cream sauce} 0.002532996 0.4318182 0.005865885 22.650826 19
[3] {mushroom cream sauce, pasta} => {escalope} 0.002532996 0.9500000 0.002666311 11.976387 19
[4] {parmesan cheese, tomatoes} => {frozen vegetables} 0.002133049 0.6666667 0.003199573 6.993939 16
[5] {mineral water, whole wheat pasta} => {olive oil} 0.003866151 0.4027778 0.009598720 6.115863 29
> |

```

(c)

first Generating a transaction rule for second case with maximum length of 2

```

300
301 transaction_rules2 <- apriori(d, parameter = list(support = 0.002, confidence = 0.20, maxlen = 2))
302 inspect(transaction_rules2[1:5])
303

```

```
> inspect(transaction_rules2[1:5])
```

	lhs	rhs	support	confidence	coverage	lift	count
[1]	{}	=> {mineral water}	0.238368218	0.2383682	1.000000000	1.0000000	1788
[2]	{asparagus}	=> {mineral water}	0.002133049	0.4444444	0.004799360	1.8645290	16
[3]	{candy bars}	=> {mineral water}	0.002266364	0.2328767	0.009732036	0.9769621	17
[4]	{shallot}	=> {green tea}	0.002266364	0.2931034	0.007732302	2.2185358	17
[5]	{shallot}	=> {french fries}	0.002666311	0.3448276	0.007732302	2.0175910	20

```
> |
```

here is the two rules t1 is the rule from QII-b with the greatest lift, t2 is rule with the highest lift rule for maximum length of 2

```
305 t1<-inspect(top.lift[1])
306 t2<-inspect(transaction_rules2[1])
307 |
```

```
> t1<-inspect(top.lift[1])
```

	lhs	rhs	support	confidence	coverage	lift	count
[1]	{escalope, mushroom cream sauce}	=> {pasta}	0.002532996	0.4418605	0.005732569	28.0881	19

```
> t2<-inspect(transaction_rules2[1])
```

	lhs	rhs	support	confidence	coverage	lift	count
[1]	{}	=> {mineral water}	0.2383682	0.2383682	1	1	1788

```
> |
```

- (i) The **first** one has a **better left** than the second one, The **second** one has a **better support** than the first one
- (ii)
- (iii) Selecting the rule with **highest lift** because Lift is a relative strength indicator showing the association between ,so I will select

```
> t1<-inspect(top.lift[1])
```

	lhs	rhs	support	confidence	coverage	lift	count
[1]	{escalope, mushroom cream sauce}	=> {pasta}	0.002532996	0.4418605	0.005732569	28.0881	19